

SKILLS

AI & ML	LLMs, LangChain, LangGraph, CrewAI, FAISS, Weaviate, FastAPI, MLflow, LangSmith, MCP, NumPy, Pandas
LLM & Cloud AI	OpenAI GPT-4o, Azure OpenAI, RAG Pipelines, Agentic AI, Prompt Engineering, LLM Evaluation
Programming & DB	Python, SQL
Data Engineering	Apache Spark, PySpark, Databricks, Apache Kafka, Docker, AWS Athena & S3
Tools & BI	Git, Power BI, Figma, Ollama, SQLite

TECHNICAL EXPERIENCE

AI Analytics RAG System	Nexturn
Cloud Engineer, AI/ML	January 2026 – Present

- Developed an **agentic** multi-agent RAG pipeline using **LangChain** and **LangGraph** to convert natural language queries into SQL analytics against **AWS Athena**, with **FAISS** and **Weaviate** (local & cloud) for schema retrieval.
- Engineered a production LLM evaluation pipeline using ground-truth datasets, scoring outputs on **faithfulness**, **context precision**, and **hallucination rate** via cosine similarity and semantic scoring.
- Delivered a real-time AI analytics dashboard using **React**, **Plotly**, and **SSE streaming** with AI-generated summaries and automated **PDF report exports**.

AI-Powered Recruitment Management System	Nexturn
Cloud Engineer, AI/ML	January 2026 – Present

- Built a full-stack recruitment platform with **FastAPI** and **React/Vite** frontend, integrating **OpenAI GPT-4o-mini** and **Azure OpenAI** for automated resume parsing, prescreening, and AI-based candidate scoring against job descriptions.
- Designed a multi-stage resume ingestion pipeline with **PII masking**, text extraction, and **SharePoint** integration via Microsoft Graph API for candidate pool management.
- Developed an interview scheduling module sending role-specific HTML emails with **.ics** attachments via **SendGrid API**; engineered gap detection engine reducing false-positive screening flags.

Data & Machine Learning Engineer Intern	Nexturn
February 2025 – December 2025	Hyderabad, Telangana

- Orchestrated multi-agent workflows using **CrewAI** for task decomposition; built and evaluated ML/DL models with iterative validation and performance benchmarking using **MLflow**.
- Music Genre Analysis & Prediction:** Designed an end-to-end pipeline with **Apache Kafka** for real-time ingestion and **PySpark** for distributed processing; applied **Random Forest** (One-vs-Rest) for multi-class classification; visualized results in **Power BI**.

PROJECTS

LLM Benchmark Dashboard
- Built a full-stack LLM evaluation platform benchmarking 5 local models (Gemma 2B, Phi-3 Mini, DeepSeek Coder 6.7B, Mistral 7B, Llama 3.1 8B) via Ollama and Docker — zero API cost, full offline capability. - Implemented automated evaluation scoring relevancy, faithfulness, and hallucination rate using a11-MiniLM-L6-v2 embeddings; integrated LangSmith @traceable decorators for full per-model observability across every benchmark run. - Delivered a React/TypeScript dashboard with real-time SSE progress tracking, Recharts visualizations, and a FastAPI + SQLite backend with Swagger-documented REST API.

EDUCATION

Master of Technology – Data Science	National Institute of Technology, Mizoram
2023 – 2025	
Bachelor of Technology	National Institute of Technology, Mizoram
2020 – 2023	
St. Paul’s Higher Secondary School	Mizoram
MBSE	2016 – 2020